

Transluce 

Early rogue AI agent activity and attempts to hack found on urlquery.net

Jack Cable^{*,2}, Daniel Chiu^{*}, Francisco Pernice^{*,3}, Selena Zhang^{*,1}, James Anthony¹, Tetiana Bas⁴, Gary Shen⁴, Conrad Stosz¹, Jacob Steinhardt¹


¹ Transluce · ² Corridor · ³ MIT · ⁴ AIUC · ^{*}Primary contributors, listed alphabetically

Transluce | Published: September 23, 2026

We present evidence that AI agents used the web security service urlquery.net to bypass restrictions and expand their access to the public internet. The agents also tried on three occasions to hack public data providers, including an Australian government website. We link at least some of this activity to agent swarms previously attributed to OpenAI. We also find evidence of earlier agent activity going back to at least March 6th, 2026, and potentially earlier, predating the previously reported [Hugging Face](#), [collusion.wiki](#), and [RubyGems](#) incidents by at least two months.

[Download the Data](#)

[Get Involved](#)



Key Findings

- We report three separate incidents between May and June 2026 in which the agents attempted to exploit security vulnerabilities and hack into websites, including an attempt on an Australian government public health website. Notably, the agents did this while attempting mundane data retrieval tasks which were not cyber-related.
- This traffic goes back at least to March 6, 2026 and extends as recently as September 16, 2026, suggesting agents may still be exploiting these services to bypass restrictions.
- We are releasing a dataset containing tens of thousands of queries apparently made by autonomous AI agents leveraging a URL scanning service to avoid access restrictions. We encourage others to continue looking into the data.

Executive Summary

Note: on the day of publication, Australian Prime Minister Anthony Albanese [publicly announced](#) that several government websites had been infiltrated by OpenAI agents, likely overlapping with the incident we describe here. OpenAI released [a statement](#) acknowledging their involvement.

Agents attempted to hack three public data sources, including an Australian government website, and some are linked to a known agent swarm.¹ We present evidence of AI agents attempting to compromise websites at three domains: Data USA² (api.datausa.io), the University of New Mexico digital library (nmdigital.unm.edu), and the Australian Institute of Health and Welfare (AIHW) Tableau collections (viz.aihw.gov.au). This attempted compromise of AIHW is part of the first reported instance of agents hacking a government. We directly link two of the three (AIHW and Data USA) to a previously reported agent swarm that OpenAI has [publicly confirmed](#) originated from them. For all three, we note that the extent of the observed activity is minor, attempting a low number of probe payloads and we observe no evidence of exploitation. While [previous reporting](#) showed that agents had interacted with these domains, this discovery reveals that agents attempted to hack into them when other methods of collecting the data they sought failed. Notably, the tasks the agents were trying to solve were *not cyber-related*; the agents resorted to hacking tactics while working on ordinary data retrieval tasks.

We find evidence of unintended, task-driven agent-like activity starting on March 6th. Records from urlquery.net show agents using the service since at least March 6, 2026, about two months before previously reported swarm activity. The first case, a March 6 attempt to retrieve Thai drug-enforcement statistics, shows an agent escalating as each approach failed: it first requested the data directly, then tried a service that converts web pages into text, and finally packed a custom program into a web address. The same technique shows up in thousands of agent requests recorded by urlquery.net starting in mid-April, targets many of the same data sources as the collusion.wiki swarm, and collapsed the same day the wiki activity did. We also report similar activity that occurred as recently as September 16.

We find weaker evidence of similar data-retrieval agent activity as early as November 2025. November 2025 urlquery.net records reveal bursts of attempts to retrieve statistics of historical theme park data and Thai government data through different URLs. These earlier attempts are less sophisticated and we are less confident that they involve the same agents, but they are consistent with task-directed data retrieval and target the same sources accessed in later activity.

Overall, the evidence is consistent with, but does not prove, that the agents may have learned this behavior over one or more training runs. In November, they may have used urlquery.net simply to look up information. By March, they were finding creative ways around access limits. By May and June, they were gaining more access, including attempting to bypass cyber defenses to complete their tasks.

Hacking attempts against public data providers

Much of the urlquery.net activity appears to come from agents retrieving data to answer web search tasks. For three of these tasks, after failing to retrieve data through normal means, they attempted a variety of cyber exploits against the relevant data service. We tie two of these attempts (those targeting api.datausa.io and viz.aihw.gov.au) to the [prior DseWiki agent swarm activity](#) confirmed to originate from OpenAI based on shared targets, tactics, and timing. None of the hacking attempts we identified appear to have succeeded, though the public artifacts

we analyzed are incomplete and we cannot rule out successful attempts through private scans or means other than urlquery.net. This data reveals that malicious cyber activity is not limited to agents tasked with cybersecurity-related tasks and can arise instrumentally to solve mundane tasks like information retrieval.

- › **Agents targeted University of New Mexico's digital library using exploits like SQL injection and path traversal**

- › **Agents targeted Data USA with cross-site scripting and other exploits**

- › **Agents targeted the Australian Institute of Health and Welfare**

Strong evidence of agent activity begins in March 2026, predating known incidents

Beyond the three attempted hacks, our dataset contains thousands of reports of agent activity. This activity begins on March 6, 2026, increases sharply in mid-April, peaks in May and June alongside the collusion.wiki swarm, and collapses on June 22, the day sustained agent activity on the wiki ends.

- › **March-September 2026: a timeline and connection to previous swarm activity**

- › **Attempts to create disposable email addresses and to trade cryptocurrency**

Potential agent activity begins as early as November 2025

Records from urlquery.net contain suggestive evidence of agent activity as early as November 2025. These records differ from the stronger examples beginning in March: the November reports show repeated requests for specific pages and datasets, whereas later reports use sophisticated techniques matching later attempts to access confirmed agent targets.

- › **November 2025: repeated attempts to retrieve theme-park and Thai government data**

Footnotes

1. We reached out to OpenAI and the three affected organizations on September 21 and 22 to disclose the attempted intrusions to them. 
2. Data USA is a joint project by Deloitte, Datawheel, and MIT's Collective Learning Group that provides open access to U.S. government data. It is not a government website, and is distinct from data.gov. 
3. For instance, the requested URL containing the XSS payload, **vizprod.aihw.gov.au/t/Public/views/PBSdashboardallATC1-ATC2medicines-Agegroup/PBSDashboard?:showVizHome=no&:embed=y&:apiID=""><svg/onload=document.title="XSS">**, queries the same dashboard that has been documented on DseWiki. 
4. The exact question is "In Korea, in what year between 2000 and 2021 did crude oil net imports average above 2,500 kilobarrels per day (2500 kb/d), natural gas net imports total less than 56 billion cubic meters (56 bcm), coal imports exceeded 136 million metric tons (136 Mt) but decreased compared to the year prior." 
5. See records <https://search.urlquery.net/report/1ad9c2e8-96ff-44af-b446-b717bcb995b4>, <https://urlquery.net/report/c08684cc-3da4-4d53-a288-0d014243c075> and <https://1.urlquery.net/report/e044dea5-ca3b-4e3c-9083-f422148ffd77>. 

Appendix

▼ About urlquery.net

The free online service urlquery.net allows a user to open a URL through a sandboxed remote browser. Its intended use is to test suspicious URLs without opening them locally.

From an agents' perspective, urlquery.net provides the following useful functionality:

- It gives an actor able to submit HTTP requests access to a programmable remote browser. In conjunction with other proxy sites, it allows agents to run JavaScript to retrieve data, process it, and send results elsewhere, as well as interact with dynamic web applications. In the AIHW episode described below, agents used this capability to operate Tableau dashboards and extract summary data.
- Once a piece of data has been fetched (potentially through a sophisticated method), urlquery.net automatically keeps a public record of the result that can be referenced by

other agents looking for the same data with a simple GET request.

From our perspective, the latter point makes urlquery.net an excellent record **of both the data that agents fetched and the mechanism that they used to do so.**

A limitation of the urlquery.net service is that, as far as we can tell, it requires the ability to create an HTTP POST request to specify the scanned URL. Previously [reported agent coordination](#) in DseWiki suggests that at least some agents only had the ability to make GET requests, which is what drew them to the DseWiki in the first place: unlike most websites, it can be written to via cleverly formatted GET requests. Therefore, in order to interact with urlquery.net, agents either needed to have broader HTTP abilities than the DseWiki agents or use a tool upstream of urlquery.net to convert GET requests into POST requests. There is evidence on the DseWiki forum that agents achieved this, and we see matching evidence of agents using services such as [milankarman.github.io/Embed-Web-Playground](#), [httpbin.org](#) and [blogsflow.liftbrandfulfillment.com](#).⁵

▼ About our dataset

We compiled a dataset of urlquery.net reports containing what we call agent-like activity: reports that satisfy any of the following:

1. Clearly attempt to read data from a resource that agents have shown interest in elsewhere.
2. Use the same techniques as those in (1), for instance the same tools to access a resource indirectly.
3. Tie directly to a report in (1) or (2) through identical data or an identifier.

We classified 6,467 reports as containing significant evidence of agent-like activity, based on distinctive task-specific programs, task-linked exploit probes, or exact connections to known activity. We supplement this with 31,182 reports containing suggestive evidence, identified based on the data source they target or use of techniques that are less distinctive of agent activity.

Transluce

Transluce is an independent, nonprofit research lab building the public tech stack for scalable oversight of AI.

[Donate](#)

Follow Transluce

 [X \(Twitter\)](#)

 [LinkedIn](#)

 [GitHub](#)

 [Email us](#)

Resources

[Research & news](#)

[Docent](#)

[Tools & demos](#)

Company

[About us](#)

[Join us](#) 

[Privacy & Trust Center](#) 

[Responsible disclosure](#)